

Medical Insurance Cost Prediction Project

Bert

2022-04-14

```
library(readr)
library(corrplot)

## corrplot 0.92 loaded

library(leaps)
library(glmnet)

## Warning: package 'glmnet' was built under R version 4.1.3

## Loading required package: Matrix

## Loaded glmnet 4.1-3

#import data
med = read.csv("C:/Users/bert0/Documents/Bert-school/Columbia
University/Spring 2022/PM/insurance.csv", header=TRUE)
df = read.csv("C:/Users/bert0/Documents/Bert-school/Columbia
University/Spring 2022/PM/insurance.csv", header=TRUE)

#check missing data
which(is.na(med))

## integer(0)

#change data structure
med$sex <- as.factor(med$sex)
med$smoker <- as.factor(med$smoker)
med$region <- as.factor(med$region)


levels(med$sex) <- c("female", "male")
levels(med$smoker) <- c("No", "Yes")
levels(med$region) <- c("northeast", "northwest", "southwest", "southeast")

#data visualization
summary(med)
```


##	age	sex	bmi	children	smoker
##	Min. :18.00	female:662	Min. :15.96	Min. :0.000	No :1064
##	1st Qu.:27.00	male :676	1st Qu.:26.30	1st Qu.:0.000	Yes: 274
##	Median :39.00		Median :30.40	Median :1.000	
##	Mean :39.21		Mean :30.66	Mean :1.095	
##	3rd Qu.:51.00		3rd Qu.:34.69	3rd Qu.:2.000	
##	Max. :64.00		Max. :53.13	Max. :5.000	

```
##      region      charges
## northeast:324  Min.   : 1122
## northwest:325  1st Qu.: 4740
## southwest:364  Median : 9382
## southeast:325  Mean    :13270
##                3rd Qu.:16640
##                Max.    :63770
```

```
plot(med$charges)
```




```
plot(log(med$charges))
```




```
par(mfrow = c(2,4))  
#4 numerical data  
hist(med$charges,xlab = "Insurance Cost (Y) ", main="Medical Insurance Cost  
Data")  
hist(med$age,xlab = "Age", main="Medical Insurance Cost Data")  
hist(med$bmi,xlab = "BMI", main="Medical Insurance Cost Data")  
hist(med$children,xlab = "Number of children", main="Medical Insurance Cost  
Data")  
#3 categorical data  
barplot(table(med$sex),xlab="Sex",  
main = "Medical Insurance Cost Data")  
barplot(table(med$smoker),xlab="Smoker",  
main = "Medical Insurance Cost Data")  
barplot(table(med$region),xlab="Region",  
main = "Medical Insurance Cost Data")
```

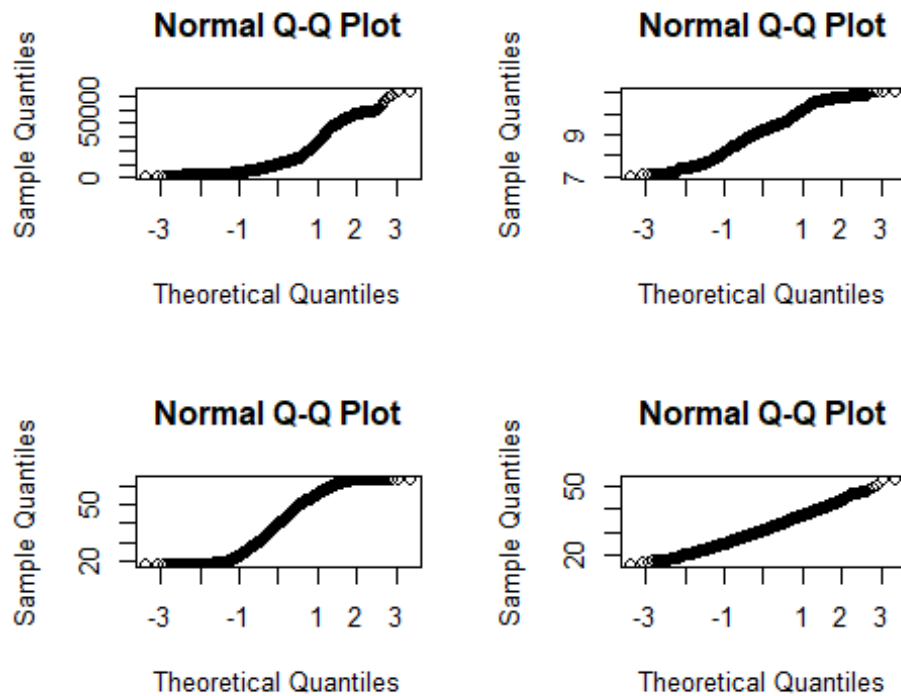
Medical Insurance Charges by Insurance Company



Medical Insurance Charges by Insurance Company



```
par(mfrow = c(2,2))
qqnorm(med$charges)
qqnorm(log(med$charges))
qqnorm(med$age)
qqnorm(med$bmi)
```



```
#Charts :1.charges, 2.ln(charges) 3.age 4.bmi
```

```
#Change charges into ln()
```

```
med$charges<-log(med$charges)
```

```
#full regression model
```

```
y=med$charges
```

```
res=lm(y~age+sex+bmi+children+smoker+region,data=med)
```

```
summary(res)
```

```
##
```

```
## Call:
```

```
## lm(formula = y ~ age + sex + bmi + children + smoker + region,
##     data = med)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -1.07186 -0.19835 -0.04917  0.06598  2.16636
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   7.0305581  0.0723960  97.112  < 2e-16 ***
## age           0.0345816  0.0008721  39.655  < 2e-16 ***
## sexmale      -0.0754164  0.0244012  -3.091  0.002038 **
## bmi           0.0133748  0.0020960   6.381  2.42e-10 ***
## children      0.1018568  0.0100995  10.085  < 2e-16 ***
## smokerYes     1.5543228  0.0302795  51.333  < 2e-16 ***
```

```
## regionnorthwest -0.0637876  0.0349057  -1.827  0.067860 .
## regionsouthwest -0.1571967  0.0350828  -4.481  8.08e-06 ***
## regionsoutheast -0.1289522  0.0350271  -3.681  0.000241 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4443 on 1329 degrees of freedom
## Multiple R-squared:  0.7679, Adjusted R-squared:  0.7666
## F-statistic: 549.8 on 8 and 1329 DF,  p-value: < 2.2e-16
```

We are going to select models, we use MSE and other values to compare models and following part have three methods to measure MSE: #1.using all data 2.validation set 3.Cross Validation

```
reg.best=regsubsets(charges~ . ,data=med,nvmax=8)
summary(reg.best)

## Subset selection object
## Call: regsubsets.formula(charges ~ ., data = med, nvmax = 8)
## 8 Variables (and intercept)
##              Forced in Forced out
## age                FALSE      FALSE
## sexmale             FALSE      FALSE
## bmi                 FALSE      FALSE
## children            FALSE      FALSE
## smokerYes          FALSE      FALSE
## regionnorthwest    FALSE      FALSE
## regionsouthwest    FALSE      FALSE
## regionsoutheast    FALSE      FALSE
## 1 subsets of each size up to 8
## Selection Algorithm: exhaustive
##      age sexmale bmi children smokerYes regionnorthwest
regionsouthwest
## 1 ( 1 ) " " " " " " " " " " " "
## 2 ( 1 ) "*" " " " " " " " " " "
## 3 ( 1 ) "*" " " " " "*" " " " " "
## 4 ( 1 ) "*" " " "*" "*" " " " " " "
## 5 ( 1 ) "*" " " "*" "*" " " " " "*"
## 6 ( 1 ) "*" " " "*" "*" " " " " "*"
## 7 ( 1 ) "*" "*" "*" "*" " " " " "*"
## 8 ( 1 ) "*" "*" "*" "*" "*" " " " "*"
##      regionsoutheast
## 1 ( 1 ) " "
## 2 ( 1 ) " "
## 3 ( 1 ) " "
## 4 ( 1 ) " "
## 5 ( 1 ) " "
## 6 ( 1 ) "*"
## 7 ( 1 ) "*"
## 8 ( 1 ) "*"

```

```
#using whole data to train
```

```
reg.Med.summary=summary(reg.best)
```

```
reg.Med.summary$outmat
```

```
##          age sexmale bmi children smokerYes regionnorthwest
regionsouthwest
## 1 ( 1 ) " " " "      " " " "      "*"      " "      " "
## 2 ( 1 ) "*" " "      " " " "      "*"      " "      " "
## 3 ( 1 ) "*" " "      " " "*"      "*"      " "      " "
## 4 ( 1 ) "*" " "      "*" "*"      "*"      " "      " "
## 5 ( 1 ) "*" " "      "*" "*"      "*"      " "      "*"
## 6 ( 1 ) "*" " "      "*" "*"      "*"      " "      "*"
## 7 ( 1 ) "*" "*"      "*" "*"      "*"      " "      "*"
## 8 ( 1 ) "*" "*"      "*" "*"      "*"      "*"      "*"
##          regionsoutheast
## 1 ( 1 ) " "
## 2 ( 1 ) " "
## 3 ( 1 ) " "
## 4 ( 1 ) " "
## 5 ( 1 ) " "
## 6 ( 1 ) "*"
## 7 ( 1 ) "*"
## 8 ( 1 ) "*"

```

```
r2=reg.Med.summary$rsq
```

```
adjr2=reg.Med.summary$adjr2
```

```
cp=reg.Med.summary$cp
```

```
bic=reg.Med.summary$bic
```

```
table_best=data.frame(model=c(1:8),r2,adjr2,cp,bic)
```

```
table_best
```

```
##  model      r2      adjr2      cp      bic
## 1      1 0.4428978 0.4424809 1856.61244 -768.341
## 2      2 0.7395465 0.7391564  159.65828 -1778.456
## 3      3 0.7572654 0.7567195   60.17950 -1865.527
## 4      4 0.7621566 0.7614429   34.16713 -1885.564
## 5      5 0.7639274 0.7630413   26.02496 -1888.365
## 6      6 0.7657049 0.7646487   17.84507 -1891.278
## 7      7 0.7673647 0.7661403   10.33949 -1893.591
## 8      8 0.7679478 0.7665509    9.00000 -1889.750

```

```
#Looking for coef
```

```
coef(reg.best,3)
```

```
## (Intercept)      age      children      smokerYes
## 7.28772342  0.03528491  0.10163109  1.54427238

```

```
coef(reg.best,4)
```

```
## (Intercept)      age      bmi      children      smokerYes
## 6.98277656  0.03478256  0.01060965  0.10119760  1.54324382


```

prefer 3 or 4 model

```
#Validation Set
set.seed(1)
train=sample(c(TRUE,FALSE), nrow(med),rep=TRUE)
#sum(train==TRUE);sum(train==FALSE)
test=(!train)
regfit.full=regsubsets(charges~.,data=med[train,],nvmax=8)
#summary(regfit.full)
test.mat=model.matrix(charges~.,data=med[test,])
val.errors=rep(NA,8)
for (i in 1:8)
{
  coefi=coef(regfit.full,id=i)
  pred=test.mat[,names(coefi)]%*%coefi
  val.errors[i]=mean((med$charges[test]-pred)^2)
}
data.frame(model=c(1:8),val.errors)

##   model val.errors
## 1      1  0.4586110
## 2      2  0.2126700
## 3      3  0.1962058
## 4      4  0.1916911
## 5      5  0.1911752
## 6      6  0.1905866
## 7      7  0.1886521
## 8      8  0.1881355

plot(val.errors)
```

```
#minimum errors' model
#coef(regfit.full, which(val.errors==min(val.errors)))
```

validation test show...

```
#prediction
predict.regsubsets=function(object,newdata,id,...)
{
  form=as.formula(object$call[[2]])
  mat=model.matrix(form,newdata)
  coefi=coef(object,id=id)
  xvars=names(coefi)
  mat[,xvars]%*%coefi
}

#Cross validation
k=10
set.seed(1)
folds=sample(1:k,nrow(med),replace=TRUE)

#folds
cv.errors=matrix(NA,k,8, dimnames=list(NULL, paste(1:8)))
for(j in 1:k){
  best.fit=regsubsets(charges~.,data=med[folds!=j,],nvmax=8)
  for(i in 1:8){
    pred=predict.regsubsets(best.fit,med[folds==j,],id=i)
    cv.errors[j,i]=mean( (med$charges[folds==j]-pred)^2)
  }
}
```


```

mean.cv.errors=apply(cv.errors,2,mean)
#draw plot
data.frame(model=c(1:8),mean.cv.errors)

##   model mean.cv.errors
## 1     1      0.4711180
## 2     2      0.2209773
## 3     3      0.2061846
## 4     4      0.2021111
## 5     5      0.2033770
## 6     6      0.2013673
## 7     7      0.1986568
## 8     8      0.1982912

plot(mean.cv.errors)

```



```

#minimum errors' model#coef(regfit.full,
which(mean.cv.errors==min(mean.cv.errors)))

```

CV shows

CONCLUSION could be: Although we have model 6 with lowest MSE, we prefer 3 factor model is the best, because it is not complicated and the accuracy is relatively high.

```


#3 factor model is the best, because it is not complicated and the accuracy
is relatively high.
res2=lm(y~age+children+smoker,data=med)
summary(res2)

```


```
##
## Call:
## lm(formula = y ~ age + children + smoker, data = med)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.94939 -0.17632 -0.04368  0.04252  2.13501
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  7.2877234   0.0387040  188.294   <2e-16 ***
## age          0.0352849   0.0008839   39.919   <2e-16 ***
## children     0.1016311   0.0102990    9.868   <2e-16 ***
## smokerYes    1.5442724   0.0307364   50.242   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4535 on 1334 degrees of freedom
## Multiple R-squared:  0.7573, Adjusted R-squared:  0.7567
## F-statistic: 1387 on 3 and 1334 DF, p-value: < 2.2e-16
```

Also, we can try to do Ridge and Lasso to reduce Variance


```
#Ridge & Lasso
# make design matrix x and response y=charges
x=model.matrix(charges~.,med)[-1]
y=med$charges
fit.ridge = glmnet(x,y,alpha=0)
plot(fit.ridge, xvar="lambda", label=TRUE)
```




```
cv.ridge = cv.glmnet(x,y,alpha =0);plot(cv.ridge)
```



```
fit.lasso = glmnet(x,y,alpha=1)
plot(fit.lasso, xvar="lambda", label=TRUE)
```



```
cv.lasso = cv.glmnet(x,y,alpha =1);plot(cv.lasso)
```



We can see which proportion of training set have best performance Both Ridge and Lasso.
Table1:Ridge;Table2:Lasso

```
#decide training set
set.seed(1)
tp=rep(0,9)
mse=rep(0,9)
a=c(0.5,0.5625,0.625,0.6875,0.75,0.8125,0.875,0.9375,1)
for (i in 1:9) {
  if (i<1) {
    train=sample(1:nrow(x), a[i]*nrow(x))
    test=(-train)
    y.test = y[-train]
    alpha0.fit =cv.glmnet(x[train,],y[train], alpha=0,
type.measure="mse",family="gaussian")
    alpha0.predicted = predict(alpha0.fit, s=alpha0.fit$lambda.1se,newx=x[-
train,])
    tp[i]=alpha0.fit$lambda.1se
    mse[i]=mean((y.test-alpha0.predicted)^2)
  }
  else {
    train=sample(1:nrow(x), a[i]*nrow(x))
    test=(train)
    y.test = y[train]
    alpha0.fit =cv.glmnet(x[train,],y[train], alpha=0,
type.measure="mse",family="gaussian")
    alpha0.predicted = predict(alpha0.fit,
```

```

s=alpha0.fit$lambda.1se,newx=x[train,])
tp[i]=alpha0.fit$lambda.1se
mse[i]=mean((y.test-alpha0.predicted)^2)
}

```

```

}

```

```

table_Ridge=data.frame(RatioOfTrainingSet = a, tp=tp,MSE=mse)
table_Ridge

```

```

##      RatioOfTrainingSet      tp      MSE
## 1          0.5000 0.1867904 0.2029582
## 2          0.5625 0.2037921 0.2294636
## 3          0.6250 0.1635435 0.1942499
## 4          0.6875 0.1896704 0.2163271
## 5          0.7500 0.1725343 0.2206320
## 6          0.8125 0.1875701 0.2154752
## 7          0.8750 0.1429752 0.2069787
## 8          0.9375 0.1517741 0.2087969
## 9          1.0000 0.1868110 0.2152947

```

#Lasso

```

for (i in 1:9) {
  if (i<1) {
    train=sample(1:nrow(x), a[i]*nrow(x))
    test=(-train)
    y.test = y[-train]
    alpha1.fit =cv.glmnet(x[train,],y[train], alpha=1,
type.measure="mse",family="gaussian")
    alpha1.predicted = predict(alpha1.fit, s=alpha1.fit$lambda.1se,newx=x[-
train,])
    tp[i]=alpha1.fit$lambda.1se
    mse[i]=mean((y.test-alpha1.predicted)^2)
  }
  else {
    train=sample(1:nrow(x), a[i]*nrow(x))
    test=(train)
    y.test = y[train]
    alpha1.fit =cv.glmnet(x[train,],y[train], alpha=1,
type.measure="mse",family="gaussian")
    alpha1.predicted = predict(alpha1.fit,
s=alpha1.fit$lambda.1se,newx=x[train,])
    tp[i]=alpha1.fit$lambda.1se
    mse[i]=mean((y.test-alpha1.predicted)^2)
  }
}

```

```

table_Lasso=data.frame(RatioOfTrainingSet = a, tp=tp,MSE=mse)
table_Lasso

```

##	RatioOfTrainingSet	tp	MSE
## 1	0.5000	0.05541572	0.2192705
## 2	0.5625	0.07042788	0.2347370
## 3	0.6250	0.04659242	0.2043661
## 4	0.6875	0.06036091	0.2075990
## 5	0.7500	0.06036840	0.2022688
## 6	0.8125	0.05063202	0.2036441
## 7	0.8750	0.04950568	0.2046174
## 8	0.9375	0.05522830	0.2120856
## 9	1.0000	0.06559293	0.2170027

conclude that ratio between 0.625-0.75 is better ratio Using Elastic_Net Regression and using 0.6875 times of data as training data

#Elastic_Net Regression

```
set.seed(1)
train=sample(1:nrow(x), 0.6875*nrow(x))
test=(-train)
y.test = y[-train]

list.of.fits = list()
for (i in 0:10) {
  fit.name = paste0("alpha",i/10)
  list.of.fits[[fit.name]] = cv.glmnet(x[train,],
y[train],alpha=i/10,type.measure="mse", family="gaussian")
}
results= data.frame()

for (i in 0:10) {
  fit.name = paste0("alpha",i/10)
  predicted =
predict(list.of.fits[[fit.name]],s=list.of.fits[[fit.name]]$lambda.1se,newx=x
[-train,])

  mse = mean((y.test - predicted)^2)
  temp = data.frame(alpha = i/10, mse=mse, fit.name = fit.name)
  results = rbind(results,temp)
}
results
```


##	alpha	mse	fit.name
## 1	0.0	0.2463103	alpha0
## 2	0.1	0.2426991	alpha0.1
## 3	0.2	0.2488905	alpha0.2
## 4	0.3	0.2554663	alpha0.3
## 5	0.4	0.2533036	alpha0.4
## 6	0.5	0.2613081	alpha0.5
## 7	0.6	0.2553572	alpha0.6
## 8	0.7	0.2593479	alpha0.7
## 9	0.8	0.2576691	alpha0.8


```
## 10 0.9 0.2514805 alpha0.9
## 11 1.0 0.2489712 alpha1

results[match(min(results$mse),results$mse),]

## alpha mse fit.name
## 2 0.1 0.2426991 alpha0.1

plot(results$alpha,results$mse)
```



```
alpha1=results[match(min(results$mse),results$mse),][1,1]
alpha1

## [1] 0.1

list.of.fits = list()
for (i in 0:10) {
  fit.name = paste0("alpha",i/100 +alpha1)
  list.of.fits[[fit.name]] = cv.glmnet(x[train,], y[train],alpha=i/100
+alpha1,type.measure="mse", family="gaussian")
}
results= data.frame()

for (i in 0:10) {
  fit.name = paste0("alpha",i/100 +alpha1)
  predicted =
predict(list.of.fits[[fit.name]],s=list.of.fits[[fit.name]]$lambda.1se,newx=x
[-train,])
```

```

mse = mean((y.test - predicted)^2)
temp = data.frame(alpha = i/100+alpha1, mse=mse, fit.name = fit.name)
results = rbind(results,temp)
}
results[match(min(results$mse),results$mse),][1,1]

## [1] 0.11

results

##      alpha      mse fit.name
## 1  0.10 0.2511796 alpha0.1
## 2  0.11 0.2430825 alpha0.11
## 3  0.12 0.2454672 alpha0.12
## 4  0.13 0.2462956 alpha0.13
## 5  0.14 0.2525605 alpha0.14
## 6  0.15 0.2511881 alpha0.15
## 7  0.16 0.2525787 alpha0.16
## 8  0.17 0.2490182 alpha0.17
## 9  0.18 0.2505073 alpha0.18
## 10 0.19 0.2541204 alpha0.19
## 11 0.20 0.2584083 alpha0.2


```

get the result that alpha=0.12 in EN regression has lowest MSE


```

x=model.matrix(charges~.,med)[-1]
y=med$charges
fit.model = glmnet(x,y,alpha=0.12)
plot(fit.model, xvar="lambda", label=TRUE)

```



```
cv.model = cv.glmnet(x,y,alpha =0.12);plot(cv.model)
```



```
#Decision Tree
```